

Textbook

- 10.2, 10.4, 10.6
- Using Weka, solve 9.1 with MLNN, SVM, and another classifier of your choice
- 11.2

10.2 Suppose that the data mining task is to cluster points (with x, y representing location) into three clusters, where the points are

$$A_1(2,10), A_2(2,5), A_3(8,4), B_1(5,8), B_2(7,5), B_3(6,4), C_1(1,2), C_2(4,9).$$

The distance function is Euclidean distance. Suppose initially we assign A_1 , B_1 , and C_1 as the center of each cluster, respectively. Use the *k-means* algorithm to show *only*

(a) The three cluster centers after the first round of execution.

```
In [1]: import numpy as np
        from sklearn.cluster import KMeans
        import warnings
        warnings.filterwarnings('ignore')
```

a)

```
In [2]: points = np.array([[2,10],[2,5],[8,4],[5,8],[7,5],[6,4],[1,2],[4,9]])
        initial_centroid = np.array([[2,10],[5,8],[1,2]])
```

```
In [3]: #Using max_iter we can check for each iteration.
        kmeans = KMeans(n_clusters=3, init=initial_centroid, max_iter=1)
        kmeans.fit(points)
```

```
Out[3]: KMeans
KMeans(init=array([[ 2, 10],
 [ 5,  8],
 [ 1,  2]]), max_iter=1,
n_clusters=3)
```

```
In [4]: kmeans.cluster_centers_
```

```
Out[4]: array([[ 2. , 10. ],
 [ 6. ,  6. ],
 [ 1.5,  3.5]])
```

The three cluster centers after the first round of execution are:

[2,10], [6,6] and [1.5,3.5]

(b) The final three clusters.

Iteration 2:

b)

```
In [10]: kmeans = KMeans(n_clusters=3, init=initial_centroid, max_iter=2)
         kmeans.fit(points)
```

```
Out[10]: KMeans
KMeans(init=array([[ 2, 10],
 [ 5,  8],
 [ 1,  2]]), max_iter=2,
n_clusters=3)
```

```
In [11]: kmeans.cluster_centers_
```

```
Out[11]: array([[3. , 9.5 ],
 [6.5 , 5.25],
 [1.5 , 3.5 ]])
```

Iteration 3:

```
In [12]: kmeans = KMeans(n_clusters=3, init=initial_centroid, max_iter=3)
kmeans.fit(points)
```

```
Out[12]: KMeans
KMeans(init=array([[ 2, 10],
 [ 5,  8],
 [ 1,  2]]), max_iter=3,
n_clusters=3)
```

```
In [13]: kmeans.cluster_centers_
```

```
Out[13]: array([[3.66666667, 9.        ],
 [7.        , 4.33333333],
 [1.5       , 3.5       ]])
```

Iteration 4:

```
In [14]: kmeans = KMeans(n_clusters=3, init=initial_centroid, max_iter=4)
kmeans.fit(points)
```

```
Out[14]: KMeans
KMeans(init=array([[ 2, 10],
 [ 5,  8],
 [ 1,  2]]), max_iter=4,
n_clusters=3)
```

```
In [15]: kmeans.cluster_centers_
```

```
Out[15]: array([[3.66666667, 9.        ],
 [7.        , 4.33333333],
 [1.5       , 3.5       ]])
```

As centers at iteration 3 and iteration 4 were same (remain unchanged), these are the final centers of the cluster.

```
In [16]: kmeans.labels_
```

```
Out[16]: array([0, 2, 1, 0, 1, 1, 2, 0])
```

Final three clusters are:

- 1) [2,10] , [5,8] , [4,9]
- 2) [8,4] , [7,5] , [6,4]
- 3) [2,5] , [1,2]

10.4 For the k-means algorithm, it is interesting to note that by choosing the initial cluster centers carefully, we may be able to not only speed up the algorithm's convergence, but also guarantee the quality of the final clustering. The k-means++ algorithm is a variant of k-means, which chooses the initial centers as follows. First, it selects one center uniformly at random from the objects in the data set. Iteratively, for each object p other than the chosen center, it chooses an object as the new center. This object is chosen at random with probability proportional to $\text{dist}(p)^2$, where $\text{dist}(p)$ is the distance from p to the closest center that has already been chosen. The iteration continues until k centers are selected.

Explain why this method will not only speed up the convergence of the k-mean algorithm, but also guarantee the quality of the final clustering results.

Answer)

Random initialization of the centroids is the major drawback of k-means algorithm. This results in incorrect and inaccurate formation of clusters. Points that are selected as initial chosen centroids have higher possibility of already existing in different clusters. This ensures decrease in the runtime and provides better quality of the clusters. Initialization of the clusters highly matters in K-means. It will speed up the convergence process as it will avoid initialization of initial clusters which are similar to each other.

10.6 Both *k-means* and *k-medoids* algorithms can perform effective clustering.

(a) Illustrate the strength and weakness of *k-means* in comparison with *k-medoids*.

Answer)

K - Means	K- Medoids
Affected by outliers (Weakness)	More robust to outliers (Strength)
Faster and efficient (Strength)	Comparatively slower (Weakness)

Means are generally more affected by outliers because average can change quickly if a value is too small or too big. While, in case of median, center value is considered and because of which outliers and noise don't have significant impact.

(b) Illustrate the strength and weakness of these schemes in comparison with a hierarchical clustering scheme (e.g., AGNES).

Answer)

Partitioning	Hierarchical
Weakness 1. Need to specify the number of clusters to partition which may be unknown information before some testing 2. Less efficient	Strength 1. No need to specify the number of clusters beforehand. 2. More efficient
Strength 1. Steps can be undone 2. Quality is good	Weakness 1. Steps can't be undone 2. Quality can be bad

Using Weka, solve 9.1 with MLNN, SVM, and another classifier of your choice

9.1 The following table consists of training data from an employee database. The data have been generalized. For example, "31 ... 35" for *age* represents the age range of 31 to 35. For a given row entry, *count* represents the number of data tuples having the values for *department*, *status*, *age*, and *salary* given in that row.

<i>department</i>	<i>status</i>	<i>age</i>	<i>salary</i>	<i>count</i>
sales	senior	31 ... 35	46K ... 50K	30
sales	junior	26 ... 30	26K ... 30K	40
sales	junior	31 ... 35	31K ... 35K	40
systems	junior	21 ... 25	46K ... 50K	20
systems	senior	31 ... 35	66K ... 70K	5
systems	junior	26 ... 30	46K ... 50K	3
systems	senior	41 ... 45	66K ... 70K	3
marketing	senior	36 ... 40	46K ... 50K	10
marketing	junior	31 ... 35	41K ... 45K	4
secretary	senior	46 ... 50	36K ... 40K	4
secretary	junior	26 ... 30	26K ... 30K	6

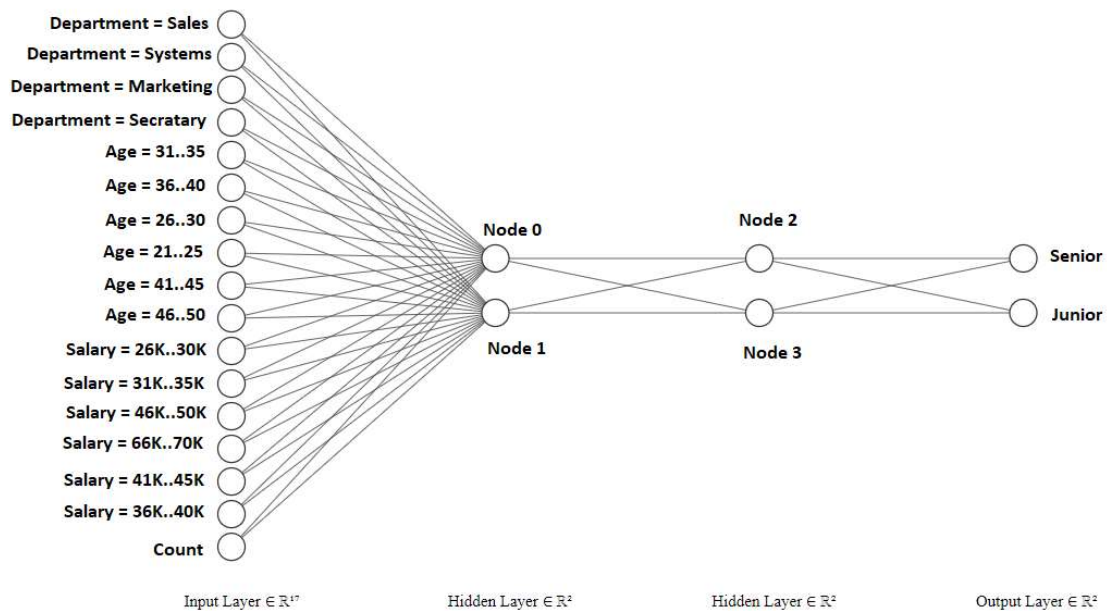
Let *status* be the class-label attribute.

- Design a multilayer feed-forward neural network for the given data. Label the nodes in the input and output layers.
- Using the multilayer feed-forward neural network obtained in (a), show the weight values after one iteration of the backpropagation algorithm, given the training instance "(sales, senior, 31 ... 35, 46K ... 50K)". Indicate your initial weight values and biases and the learning rate used.

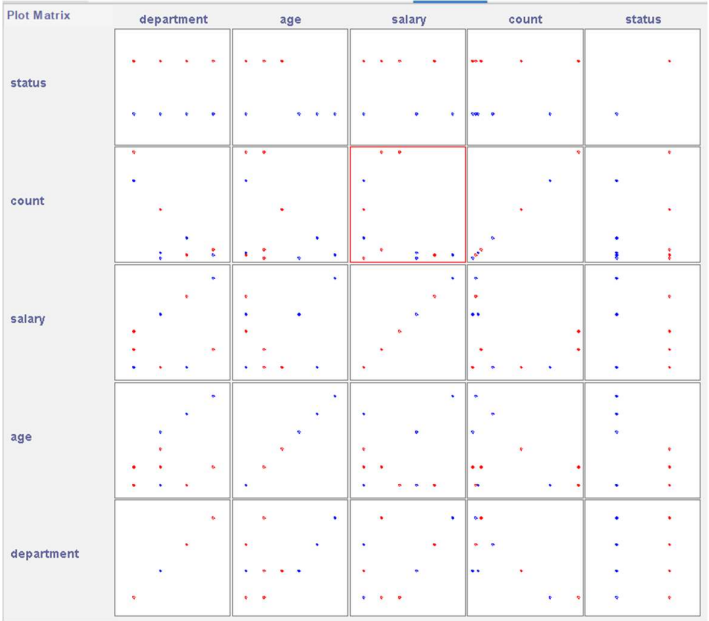



a)

Multilayer feed-forward neural network:



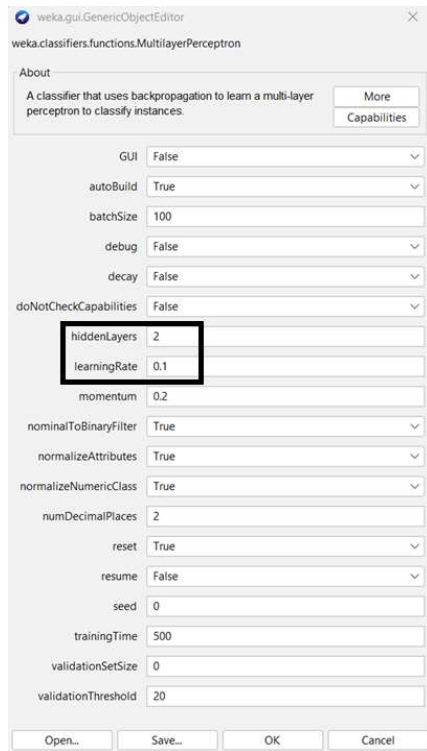
Using Weka



b)

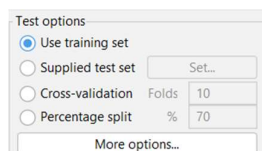
- (b) Using the multilayer feed-forward neural network obtained in (a), show the weight values after one iteration of the backpropagation algorithm, given the training instance “(sales, senior, 31 ... 35, 46K ... 50K)”. Indicate your initial weight values and biases and the learning rate used.

For multilayer feed-forward neural network, weights are:



The screenshot shows the 'weka.gui.GenericObjectEditor' window for the 'weka.classifiers.functions.MultilayerPerceptron' classifier. The 'About' section describes it as a classifier that uses backpropagation to learn a multi-layer perceptron to classify instances. The 'More' and 'Capabilities' buttons are visible. The 'GUI' checkbox is unchecked. The 'autoBuild' checkbox is checked. The 'batchSize' is set to 100. The 'debug' checkbox is unchecked. The 'decay' checkbox is unchecked. The 'doNotCheckCapabilities' checkbox is unchecked. The 'hiddenLayers' field is set to 2 and is highlighted with a black box. The 'learningRate' field is set to 0.1 and is also highlighted with a black box. The 'momentum' is set to 0.2. The 'nominalToBinaryFilter' checkbox is checked. The 'normalizeAttributes' checkbox is checked. The 'normalizeNumericClass' checkbox is checked. The 'numDecimalPlaces' is set to 2. The 'reset' checkbox is checked. The 'resume' checkbox is unchecked. The 'seed' is set to 0. The 'trainingTime' is set to 500. The 'validationSetSize' is set to 0. The 'validationThreshold' is set to 20. At the bottom, there are buttons for 'Open...', 'Save...', 'OK', and 'Cancel'.

Using training dataset.



The screenshot shows the 'Test options' dialog box. The 'Use training set' radio button is selected. The 'Set...' button is visible next to the 'Supplied test set' option. The 'Cross-validation' option has a 'Folds' field set to 10. The 'Percentage split' option has a '%' field set to 70. At the bottom, there is a 'More options...' button.

Hidden Layers = 2, Learning Rate = 0.1

Weights are:

```
Classifier output

=== Run information ===

Scheme:      weka.classifiers.functions.MultilayerPerceptron -L 0.3 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H 2
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    evaluate on training data

=== Classifier model (full training set) ===

Sigmoid Node 0
  Inputs  Weights
  Threshold 3.4851346848408755
  Node 2 -4.363442578932802
  Node 3 -3.2560705531843728
Sigmoid Node 1
  Inputs  Weights
  Threshold -3.485352476655363
  Node 2 4.343039891611943
  Node 3 3.277299190573678
Sigmoid Node 2
  Inputs  Weights
  Threshold 0.12841960905221478
  Attrib department=sales -0.5576542776440633
  Attrib department=systems 0.6331491744498269
  Attrib department=marketing 0.2422522180790835
  Attrib department=secretary -0.47183089359765656
  Attrib age=31-35 -0.4911086033732522
  Attrib age=26-30 1.769224395897942
  Attrib age=21-25 1.5078671477123209
  Attrib age=41-45 -0.7822834183690296
  Attrib age=36-40 -1.2821814881474034
  Attrib age=46-50 -0.9052525604289394
  Attrib salary=46K-50K -1.1197280374674266
  Attrib salary=26K-30K 0.603707075653472
  Attrib salary=31K-35K 1.6392077812527839
  Attrib salary=66K-70K -2.128493101438507
  Attrib salary=41K-45K 1.510905496929731
  Attrib salary=36K-40K -0.9092508064546239
  Attrib count 0.5955476674858969
Sigmoid Node 3
  Inputs  Weights
  Threshold 0.09016921625334368
  Attrib department=sales -0.4782474015342242
  Attrib department=systems 0.49698630289865714
  Attrib department=marketing 0.21278267210539267
  Attrib department=secretary -0.4287448956888747
  Attrib age=31-35 -0.39090284241848144
  Attrib age=26-30 1.5649571870867165
  Attrib age=21-25 1.2752871060028073
  Attrib age=41-45 -0.6633160123177847
  Attrib age=36-40 -1.1267201193365655
  Attrib age=46-50 -0.8057395035681446
  Attrib salary=46K-50K -0.8968887314853673
  Attrib salary=26K-30K 0.5752551319783026
  Attrib salary=31K-35K 1.3948924225197412
  Attrib salary=66K-70K -1.7755708392366605
  Attrib salary=41K-45K 1.2691247600700892
  Attrib salary=36K-40K -0.7219787887416631
  Attrib count 0.5041282077292387

Class senior
  Input
  Node 0
Class junior
  Input
  Node 1
```

```
Time taken to build model: 0.02 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

=== Summary ===

Correctly Classified Instances      11          100   %
Incorrectly Classified Instances    0           0   %
Kappa statistic                     1
Mean absolute error                 0.0298
Root mean squared error            0.0307
Relative absolute error             5.9974 %
Root relative squared error        6.1709 %
Total Number of Instances          11

=== Detailed Accuracy By Class ===

              TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
1.000  0.000  1.000  1.000  1.000  1.000  1.000  1.000  1.000  senior
1.000  0.000  1.000  1.000  1.000  1.000  1.000  1.000  1.000  junior
Weighted Avg.  1.000  0.000  1.000  1.000  1.000  1.000  1.000  1.000

=== Confusion Matrix ===

a b  <-- classified as
5 0 | a = senior
0 6 | b = junior
```


Using 70-30 split.

Test options

☐ Use training set

☐ Supplied test set

☐ Cross-validation Folds

☒ Percentage split %

Multilayer feed-forward neural network

Weights are:

```
Classifier output
=== Run information ===

Scheme:      weka.classifiers.functions.MultilayerPerceptron -L 0.1 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H 2
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    split 70.0% train, remainder test

=== Classifier model (full training set) ===

Sigmoid Node 0
  Inputs  Weights
  Threshold 2.7718152052617615
  Node 2 -3.7612727795000502
  Node 3 -2.4756334465425955
Sigmoid Node 1
  Inputs  Weights
  Threshold -2.7716965117260437
  Node 2 3.7412409700861677
  Node 3 2.4961510271168956
Sigmoid Node 2
  Inputs  Weights
  Threshold 0.11427927065628996
  Attrib department=sales -0.4905874198192064
  Attrib department=systems 0.5674878193589538
  Attrib department=marketing 0.21676922615655067
  Attrib department=secretary -0.419472727617253
  Attrib age=31-35 -0.41904617553810164
  Attrib age=26-30 1.6317956458031824
  Attrib age=21-25 1.3684879231458835
  Attrib age=41-45 -0.7259713069235665
  Attrib age=36-40 -1.1554849005186978
  Attrib age=46-50 -0.8269547591515234
```

```
Classifier output
  Attrib age=46-50 -0.8269547591515234
  Attrib salary=46K-50K -1.005766109748428
  Attrib salary=26K-30K 0.5905861520562155
  Attrib salary=31K-35K 1.4425031362818082
  Attrib salary=66K-70K -1.9163270671276526
  Attrib salary=41K-45K 1.3728662557744244
  Attrib salary=36K-40K -0.830952605177208
  Attrib count 0.5456685767301938
Sigmoid Node 3
  Inputs  Weights
  Threshold 0.06703303302989854
  Attrib department=sales -0.3713443399723155
  Attrib department=systems 0.3969489178990156
  Attrib department=marketing 0.1714314892756246
  Attrib department=secretary -0.34798862297447886
  Attrib age=31-35 -0.29668089039795056
  Attrib age=26-30 1.352778827643681
  Attrib age=21-25 1.0618954449777438
  Attrib age=41-45 -0.557207052234878
  Attrib age=36-40 -0.9298016753817471
  Attrib age=46-50 -0.6843636932759224
  Attrib salary=46K-50K -0.7167135160186119
  Attrib salary=26K-30K 0.5446796518998308
  Attrib salary=31K-35K 1.1135001146807275
  Attrib salary=66K-70K -1.4574783006232468
  Attrib salary=41K-45K 1.053990516508938
  Attrib salary=36K-40K -0.6006029784494419
  Attrib count 0.4311599550271979
Class senior
  Input
  Node 0
Class junior
  Input
  Node 1

Time taken to build model: 0.09 seconds
```



```

=== Evaluation on test split ===

Time taken to test model on test split: 0 seconds

=== Summary ===

Correctly Classified Instances      2          66.6667 %
Incorrectly Classified Instances    1          33.3333 %
Kappa statistic                    0.4
Mean absolute error                 0.3894
Root mean squared error             0.5438
Relative absolute error             77.8766 %
Root relative squared error        108.7506 %
Total Number of Instances          3

=== Detailed Accuracy By Class ===

          TP Rate  FP Rate  Precision  Recall  F-Measure  MCC   ROC Area  PRC Area  Class
          1.000    0.500    0.500    1.000    0.667     0.500  0.500    0.500    senior
          0.500    0.000    1.000    0.500    0.667     0.500  0.500    0.833    junior
Weighted Avg.   0.667    0.167    0.833    0.667    0.667     0.500  0.500    0.722

=== Confusion Matrix ===

a b  <-- classified as
1 0 | a = senior
1 1 | b = junior

```

SVM

Using training dataset.

- Lib SVM

```

Classifier output

=== Run information ===

Scheme:      weka.classifiers.functions.LibSVM -S 0 -K 2 -D 3 -G 0.0 -R 0.0 -N 0.5 -M 40.0 -C 1.0 -E 0.001 -P 0.1 -model "C:\Program Files\Weka-3-8-6" -seed 1
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    evaluate on training data

=== Classifier model (full training set) ===

LibSVM wrapper, original code by Yasser EL-Manzalawy (= WLSVM)

Time taken to build model: 0 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

=== Summary ===

Correctly Classified Instances      11          100 %
Incorrectly Classified Instances    0           0 %
Kappa statistic                    1
Mean absolute error                 0
Root mean squared error             0
Relative absolute error             0 %
Root relative squared error         0 %
Total Number of Instances          11

=== Detailed Accuracy By Class ===

          TP Rate  FP Rate  Precision  Recall  F-Measure  MCC   ROC Area  PRC Area  Class
          1.000    0.000    1.000    1.000    1.000     1.000  1.000    1.000    senior
          1.000    0.000    1.000    1.000    1.000     1.000  1.000    1.000    junior
Weighted Avg.   1.000    0.000    1.000    1.000    1.000     1.000  1.000    1.000

=== Confusion Matrix ===

a b  <-- classified as
5 0 | a = senior
0 6 | b = junior

```

- SMO

```
Classifier output

=== Run information ===

Scheme:      weka.classifiers.functions.SMO -C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K "weka.classifiers.functions.supportVector.PolyKernel -E 1.0 -C 250007" -calibrator "w
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    evaluate on training data

=== Classifier model (full training set) ===

SMO

Kernel used:
  Linear Kernel:  $K(x,y) = \langle x,y \rangle$ 

Classifier for classes: senior, junior

BinarySMO

Machine linear: showing attribute weights, not support vectors.

-0.1307 * (normalized) department=sales
+ 0.3128 * (normalized) department=systems
+ 0.0729 * (normalized) department=marketing
+ -0.255 * (normalized) department=secretary
+ -0.1664 * (normalized) age=31-35
+ 0.9261 * (normalized) age=26-30
+ 0.7862 * (normalized) age=21-25
+ -0.2169 * (normalized) age=41-45
+ -0.8913 * (normalized) age=36-40
+ -0.4376 * (normalized) age=46-50
+ -0.3616 * (normalized) salary=46K-50K

+ 0.1826 * (normalized) salary=26K-30K
+ 0.8693 * (normalized) salary=31K-35K
+ -1.2169 * (normalized) salary=66K-70K
+ 0.9643 * (normalized) salary=41K-45K
+ -0.4376 * (normalized) salary=36K-40K
+ 0.3071 * (normalized) count
+ 0.1217

Number of kernel evaluations: 61 (92.269% cached)

Time taken to build model: 0.01 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

=== Summary ===

Correctly Classified Instances      11      100    %
Incorrectly Classified Instances    0        0    %
Kappa statistic                     1
Mean absolute error                 0
Root mean squared error             0
Relative absolute error             0    %
Root relative squared error         0    %
Total Number of Instances          11

=== Detailed Accuracy By Class ===

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
      1.000    0.000    1.000     1.000    1.000     1.000    1.000    1.000    senior
      1.000    0.000    1.000     1.000    1.000     1.000    1.000    1.000    junior
Weighted Avg.  1.000    0.000    1.000     1.000    1.000     1.000    1.000    1.000

=== Confusion Matrix ===

a b  <-- classified as
5 0 | a = senior
0 6 | b = junior
```

Using 70-30 split.

- SMO

Classifier output

```
=== Run information ===

Scheme:      weka.classifiers.functions.SMO -C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K "weka.classifiers.functions.supportVector.PolyKernel -E 1.0 -C 250007" -calibrator "w
Relation:     newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    split 70.0% train, remainder test

=== Classifier model (full training set) ===

SMO

Kernel used:
  Linear Kernel: K(x,y) = <x,y>

Classifier for classes: senior, junior

BinarySMO

Machine linear: showing attribute weights, not support vectors.

-0.1307 * (normalized) department=sales
+ 0.3128 * (normalized) department=systems
+ 0.0729 * (normalized) department=marketing
+ -0.255 * (normalized) department=secretary
+ -0.1664 * (normalized) age=31-35
+ 0.9261 * (normalized) age=26-30
+ 0.7862 * (normalized) age=21-25
+ -0.2169 * (normalized) age=41-45
+ -0.8913 * (normalized) age=36-40
+ -0.4376 * (normalized) age=46-50
+ -0.3616 * (normalized) salary=46K-50K
```

Classifier output

```
+ -0.3616 * (normalized) salary=46K-50K
+ 0.1826 * (normalized) salary=26K-30K
+ 0.8693 * (normalized) salary=31K-35K
+ -1.2169 * (normalized) salary=66K-70K
+ 0.9643 * (normalized) salary=41K-45K
+ -0.4376 * (normalized) salary=36K-40K
+ 0.3071 * (normalized) count
+ 0.1217

Number of kernel evaluations: 61 (92.269% cached)

Time taken to build model: 0.02 seconds

=== Evaluation on test split ===

Time taken to test model on test split: 0 seconds

=== Summary ===

Correctly Classified Instances      2          66.6667 %
Incorrectly Classified Instances    1          33.3333 %
Kappa statistic                    0.4
Mean absolute error                 0.3333
Root mean squared error            0.5774
Relative absolute error             66.6667 %
Root relative squared error        115.4701 %
Total Number of Instances          3

=== Detailed Accuracy By Class ===

              TP Rate  FP Rate  Precision  Recall   F-Measure  MDC       ROC Area  PRC Area  Class
1.000    0.500    0.500    1.000    0.667    0.500    0.750    0.500    senior
0.500    0.000    1.000    0.500    0.667    0.500    0.750    0.833    junior
Weighted Avg.    0.667    0.167    0.833    0.667    0.667    0.500    0.750    0.722
```

Confusion Matrix

```
a b  <-- classified as
1 0 | a = senior
1 1 | b = junior
```

- **Lib SVM**

```

Classifier output

=== Run information ===

Scheme:      weka.classifiers.functions.LibSVM -S 0 -K 2 -D 3 -G 0.0 -R 0.0 -N 0.5 -M 40.0 -C 1.0 -E 0.001 -P 0.1 -model "C:\\Program Files\\Weka-3-8-6" -seed 1
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    split 70.0% train, remainder test

=== Classifier model (full training set) ===

LibSVM wrapper, original code by Yasser EL-Manzalawy (= WLSVM)

Time taken to build model: 0.01 seconds

=== Evaluation on test split ===

Time taken to test model on test split: 0 seconds

=== Summary ===

Correctly Classified Instances      2           66.6667 %
Incorrectly Classified Instances    1           33.3333 %
Kappa statistic                    0.4
Mean absolute error                 0.3333
Root mean squared error            0.5774
Relative absolute error            66.6667 %
Root relative squared error       115.4701 %
Total Number of Instances          3

=== Detailed Accuracy By Class ===

              TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
              1.000    0.500    0.500     1.000    0.667     0.500    0.750    0.500    senior
              0.500    0.000    1.000     0.500    0.667     0.500    0.750    0.833    junior
Weighted Avg.   0.667    0.167    0.833     0.667    0.667     0.500    0.750    0.722

=== Confusion Matrix ===

a b   <-- classified as
1 0 | a = senior
1 1 | b = junior

```

I will be using Logistic Regression classifier.

Logistic Regression

Using training dataset.

```

Classifier output

=== Run information ===

Scheme:      weka.classifiers.functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4
Relation:    newdata
Instances:    11
Attributes:   5
              department
              age
              salary
              count
              status
Test mode:    evaluate on training data

=== Classifier model (full training set) ===

Logistic Regression with ridge parameter of 1.0E-8
Coefficients...

Variable      Class
senior
=====
department=sales      15.4121
department=systems    -11.7066
department=marketing  -7.3458
department=secretary   5.0068
age=31-35             11.1206
age=26-30             -19.7511
age=21-25             -31.6926
age=41-45             14.3942
age=36-40             19.1757
age=46-50             14.3878
salary=46K-50K        10.7576
salary=26K-30K        -5.9401
salary=31K-35K        -37.0433
salary=66K-70K        19.7914
salary=41K-45K        -32.3982
salary=36K-40K        14.3878
count              0.0727

```

```

Classifier output
Count 0.0727
Intercept -6.1169

Odds Ratios...
Variable Class
=====
department=sales 4935937.8646
department=systems 0
department=marketing 0.0006
department=secretary 149.4213
age=31-35 67546.1661
age=26-30 0
age=21-25 0
age=41-45 1783732.7191
age=36-40 212756728.9276
age=46-50 1772393.1928
salary=46K-50K 46983.4626
salary=26K-30K 0.0026
salary=31K-35K 0
salary=66K-70K 393836095.7563
salary=41K-45K 0
salary=36K-40K 1772393.1928
count 1.0754

Time taken to build model: 0.01 seconds

=== Evaluation on training set ===

Time taken to test model on training data: 0 seconds

```

```

=== Summary ===

Correctly Classified Instances 11 100 %
Incorrectly Classified Instances 0 0 %
Kappa statistic 1
Mean absolute error 0
Root mean squared error 0
Relative absolute error 0.0001 %
Root relative squared error 0.0001 %
Total Number of Instances 11

=== Detailed Accuracy By Class ===

TP Rate FP Rate Precision Recall F-Measure MCC ROC Area PRC Area Class
1.000 0.000 1.000 1.000 1.000 1.000 1.000 1.000 senior
1.000 0.000 1.000 1.000 1.000 1.000 1.000 1.000 junior
Weighted Avg. 1.000 0.000 1.000 1.000 1.000 1.000 1.000 1.000

=== Confusion Matrix ===

a b <-- classified as
5 0 | a = senior
0 6 | b = junior

```

Logistic Regression

Using 70-30 split

```

Classifier output

=== Run information ===

Scheme: weka.classifiers.functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4
Relation: newdata
Instances: 11
Attributes: 5
    department
    age
    salary
    count
    status
Test mode: split 70.0% train, remainder test

=== Classifier model (full training set) ===

Logistic Regression with ridge parameter of 1.0E-8
Coefficients...

Variable Class
=====
department=sales 15.4121
department=systems -11.7066
department=marketing -7.3458
department=secretary 5.0068
age=31-35 11.1206
age=26-30 -19.7511
age=21-25 -31.6926
age=41-45 14.3942
age=36-40 19.1757
age=46-50 14.3878
salary=46K-50K 10.7576
salary=26K-30K -5.9401
salary=31K-35K -37.0433
salary=66K-70K 19.7914
salary=41K-45K -32.3982
salary=36K-40K 14.3878
count 0.0727

```

```
Intercept -6.1169
```

```
Odds Ratios...
```

```
Variable Class
=====
department=sales 4935937.8646
department=systems 0
department=marketing 0.0006
department=secretary 149.4213
age=31-35 67546.1661
age=26-30 0
age=21-25 0
age=41-45 1783732.7191
age=36-40 212756728.9276
age=46-50 1772393.1928
salary=46K-50K 46983.4626
salary=26K-30K 0.0026
salary=31K-35K 0
salary=66K-70K 393836095.7563
salary=41K-45K 0
salary=36K-40K 1772393.1928
count 1.0754
```

```
Time taken to build model: 0.02 seconds
```

```
=== Evaluation on test split ===
```

```
Time taken to test model on test split: 0 seconds
```

```
=== Summary ===
```

```
Correctly Classified Instances 2 66.6667 %
Incorrectly Classified Instances 1 33.3333 %
Kappa statistic 0.4
Mean absolute error 0.3333
Root mean squared error 0.5774
Relative absolute error 66.6667 %
Root relative squared error 115.4701 %
Total Number of Instances 3
```

```
=== Detailed Accuracy By Class ===
```

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	1.000	0.500	0.500	1.000	0.667	0.500	0.500	0.500	senior
	0.500	0.000	1.000	0.500	0.667	0.500	0.500	0.833	junior
Weighted Avg.	0.667	0.167	0.833	0.667	0.667	0.500	0.500	0.722	

```
=== Confusion Matrix ===
```

```
a b <-- classified as
1 0 | a = senior
1 1 | b = junior
```

11.2 *AllElectronics* carries 1000 products, P_1, \dots, P_{1000} . Consider customers Ada, Bob, and Cathy such that Ada and Bob purchase three products in common, P_1, P_2 , and P_3 . For the other 997 products, Ada and Bob independently purchase seven of them randomly. Cathy purchases 10 products, randomly selected from the 1000 products. In Euclidean distance, what is the probability that $\text{dist. Ada, Bob} / > \text{dist. Ada, Cathy}$? What if Jaccard similarity (Chapter 2) is used? What can you learn from this example?

Answer) Codes are in Jupyter Notebook

For Ada,Bob

```
AB_df #Ada,Bob
```

	I	Probability	Euclidean Distance	Jaccard
0	3.0	9.517334e-01	3.741657	0.176471
1	4.0	4.739323e-02	3.464102	0.250000
2	5.0	8.660691e-04	3.162278	0.333333
3	6.0	7.319719e-06	2.828427	0.428571
4	7.0	2.966451e-08	2.449490	0.538462
5	8.0	5.404466e-11	2.000000	0.666667
6	9.0	3.643051e-14	1.414214	0.818182
7	10.0	5.256928e-18	0.000000	1.000000

For Ada,Cathy

```
#Ada,Cathy
AC_df
```

	J	Probability	Euclidean Distance	Jaccard
0	1.0	9.214765e-02	4.242641	0.052632
1	2.0	3.800387e-03	4.000000	0.111111
2	3.0	8.247703e-05	3.741657	0.176471
3	4.0	1.026772e-06	3.464102	0.250000
4	5.0	7.505338e-09	3.162278	0.333333
5	6.0	3.171627e-11	2.828427	0.428571
6	7.0	7.344917e-14	2.449490	0.538462
7	8.0	8.363392e-17	2.000000	0.666667
8	9.0	3.758406e-20	1.414214	0.818182
9	10.0	3.796369e-24	0.000000	1.000000

Used

$\binom{7}{i-3} * \binom{990}{10-i} / \binom{997}{7}$ for calculation probability for Ada,Bob

$\binom{10}{j} * \binom{990}{10-j} / \binom{1000}{10}$ for calculation probability for Ada,Cathy

Probability that $dist.Ada,Bob / > dist.Ada,Cathy$

Using Euclidean

```
sum_euc = 0
for i in range(2, 9):
    for j in range(i+1,10):
        sum_euc += pr_i[i]*pr_j[j]
sum_euc
```

9.847436804201952e-07

Using Jaccard

```
sum_jac = 0
for i in range(2, 10):
    for j in range(0,i-1):
        sum_jac += pr_i[i]*pr_j[j]
sum_jac
```

0.09233115445920217

Used

$\sum_{i=3}^9 (\Pr(i) \sum_{j=i+1}^{10} \Pr(j))$ for Euclidean Distance. Value is 9.847436804201952e-07

$\sum_{i=3}^{10} (\Pr(i) \sum_{j=1}^{i-1} \Pr(j))$ for Jaccard Similarity. Value is 0.09233115445920217

Higher the Jaccard similarity means similar are the two quantities, while higher the Euclidean distance, more different will be two quantities.

Also, Jaccard similarity has a range between 0 to 1. Euclidean distance can only be non-negative.