

# Data Mining Project - Sustainability Around The World

Aradhya Mathur and Ozlem Gunes

## Time Series Analysis

```
In [1]: # Import libs
import pandas as pd
import numpy as np
import os, fnmatch
import matplotlib.pyplot as plt
import statsmodels.api as sm
import sklearn.metrics as mt
import seaborn as sns
from pmdarima.arima import auto_arima
from chart_studio.plotly import plot_mpl
from statsmodels.tsa.seasonal import seasonal_decompose
import datetime
```

## Read Data

```
In [4]: data_path = "./master.csv"
```

```
In [5]: data = pd.read_csv(data_path)
```

```
In [6]: data[0:5]
```

```
Out[6]:
```

	year	iso	country	region	EPI	AIR	H2O	BDH	WRS
0	2022	AFG	Afghanistan	Southern Asia	43.6	15.5	28.1	30.7	0.0
1	2022	ALB	Albania	Eastern Europe	47.1	37.5	54.1	63.9	1.9
2	2022	DZA	Algeria	Greater Middle East	29.6	39.4	53.3	22.7	33.1
3	2022	AGO	Angola	Sub-Saharan Africa	30.5	23.1	12.8	30.1	0.0
4	2022	ATG	Antigua and Barbuda	Latin America & Caribbean	52.4	56.5	50.1	54.2	15.7

## Preprocess and Analysis

Get countries with sufficient data points

```
In [7]: selected_countries = data[data['year'].isin([2018,2020,2022])]['country'].unique()
```

```
In [8]: selected_countries
```

```
Out[8]: array(['Afghanistan', 'Albania', 'Algeria', 'Angola',
'Antigua and Barbuda', 'Argentina', 'Armenia', 'Australia',
'Austria', 'Azerbaijan', 'Bahamas', 'Bahrain', 'Bangladesh',
'Barbados', 'Belarus', 'Belgium', 'Belize', 'Benin', 'Bhutan',
'Bolivia', 'Bosnia and Herzegovina', 'Botswana', 'Brazil',
'Brunei Darussalam', 'Bulgaria', 'Burkina Faso', 'Burundi',
'Cabo Verde', 'Cambodia', 'Cameroon', 'Canada',
'Central African Republic', 'Chad', 'Chile', 'China', 'Colombia',
'Comoros', 'Costa Rica', 'Cote d'Ivoire', 'Croatia', 'Cuba',
'Cyprus', 'Czech Republic', 'Dem. Rep. Congo', 'Denmark',
'Djibouti', 'Dominica', 'Dominican Republic', 'Ecuador', 'Egypt',
'El Salvador', 'Equatorial Guinea', 'Eritrea', 'Estonia',
'Eswatini', 'Ethiopia', 'Fiji', 'Finland', 'France', 'Gabon',
'Gambia', 'Georgia', 'Germany', 'Ghana', 'Greece', 'Grenada',
'Guatemala', 'Guinea', 'Guinea-Bissau', 'Guyana', 'Haiti',
'Honduras', 'Hungary', 'Iceland', 'India', 'Indonesia', 'Iran',
'Iraq', 'Ireland', 'Israel', 'Italy', 'Jamaica', 'Japan', 'Jordan',
'Kazakhstan', 'Kenya', 'Kiribati', 'Kuwait', 'Kyrgyzstan', 'Laos',
'Latvia', 'Lebanon', 'Lesotho', 'Liberia', 'Lithuania',
'Luxembourg', 'Madagascar', 'Malawi', 'Malaysia', 'Maldives',
'Mali', 'Malta', 'Marshall Islands', 'Mauritania', 'Mauritius',
'Mexico', 'Micronesia', 'Moldova', 'Mongolia', 'Montenegro',
'Morocco', 'Mozambique', 'Myanmar', 'Namibia', 'Nepal',
'Netherlands', 'New Zealand', 'Nicaragua', 'Niger', 'Nigeria',
'North Macedonia', 'Norway', 'Oman', 'Pakistan', 'Panama',
'Papua New Guinea', 'Paraguay', 'Peru', 'Philippines', 'Poland',
'Portugal', 'Qatar', 'Republic of Congo', 'Romania', 'Russia',
'Rwanda', 'Saint Lucia', 'Saint Vincent and the Grenadines',
'Samoa', 'Sao Tome and Principe', 'Saudi Arabia', 'Senegal',
'Serbia', 'Seychelles', 'Sierra Leone', 'Singapore', 'Slovakia',
'Slovenia', 'Solomon Islands', 'South Africa', 'South Korea',
'Spain', 'Sri Lanka', 'Sudan', 'Suriname', 'Sweden', 'Switzerland',
'Taiwan', 'Tajikistan', 'Tanzania', 'Thailand', 'Timor-Leste',
'Togo', 'Tonga', 'Trinidad and Tobago', 'Tunisia', 'Turkey',
'Turkmenistan', 'Uganda', 'Ukraine', 'United Arab Emirates',
'United Kingdom', 'United States of America', 'Uruguay',
'Uzbekistan', 'Vanuatu', 'Venezuela', 'Viet Nam', 'Zambia',
'Zimbabwe', "CÃfÃ´te d'Ivoire", 'Macedonia',
'SÃfÃ©o TomÃfÃ© and PrÃfÃ\xadncipe', 'Swaziland'], dtype=object)
```

```
In [9]: data = data[data['country'].isin(selected_countries)]
```

```
In [10]: data
```

Out[10]:

	year	iso	country	region	EPI	AIR	H2O	BDH	WRS
0	2022	AFG	Afghanistan	Southern Asia	43.60	15.50	28.10	30.70	0.00
1	2022	ALB	Albania	Eastern Europe	47.10	37.50	54.10	63.90	1.90
2	2022	DZA	Algeria	Greater Middle East	29.60	39.40	53.30	22.70	33.10
3	2022	AGO	Angola	Sub-Saharan Africa	30.50	23.10	12.80	30.10	0.00
4	2022	ATG	Antigua and Barbuda	Latin America & Caribbean	52.40	56.50	50.10	54.20	15.70
...	...	...	...	...	...	...	...	...	...
1170	2010	UZB	Uzbekistan	Former Soviet States	42.27	65.07	87.41	21.75	21.53
1171	2010	VEN	Venezuela	Latin America & Caribbean	62.90	97.37	83.52	78.19	56.12
1172	2010	VNM	Viet Nam	Asia-Pacific	59.00	41.47	73.46	41.28	78.05
1173	2010	ZMB	Zambia	Sub-Saharan Africa	47.00	36.76	36.86	100.00	70.28
1174	2010	ZWE	Zimbabwe	Sub-Saharan Africa	47.82	54.81	53.32	93.75	67.69

1164 rows × 9 columns

In [11]: `data_count = data.groupby(['country'])['year'].count().reset_index(name='counts').sort_values('counts', ascending=False)`

In [12]: `selected_countries = data_count[data_count['counts'] > 6]['country'].unique()`

In [13]: `data = data[data['country'].isin(selected_countries)]`

In [14]: `data[1:5]`

Out[14]:

	year	iso	country	region	EPI	AIR	H2O	BDH	WRS
2	2022	DZA	Algeria	Greater Middle East	29.6	39.4	53.3	22.7	33.1
3	2022	AGO	Angola	Sub-Saharan Africa	30.5	23.1	12.8	30.1	0.0
5	2022	ARG	Argentina	Latin America & Caribbean	41.1	52.0	64.8	42.4	5.9
6	2022	ARM	Armenia	Former Soviet States	48.3	32.1	57.3	73.3	4.5

In [15]: `data = data.reset_index(drop=True)`

In [16]: `data.isnull().sum(axis = 0)`

Out[16]:

```

year      0
iso       0
country   0
region    0
EPI       0
AIR       0
H2O       0
BDH       0
WRS       0
dtype: int64

```

## Handle Null Values

```
In [17]: # data = data.sort_values(['country', 'year'])
# data['index'] = data.groupby(['country']).cumcount()
```

```
In [18]: idx = data[['year']].drop_duplicates().sort_values('year', axis=0).reset_index(drop=True)
idx['idx'] = idx.index
# idx['idx'] = idx['idx'].astype(str)
```

```
In [19]: idx
```

```
Out[19]:
```

	year	idx
0	2010	0
1	2012	1
2	2014	2
3	2016	3
4	2018	4
5	2020	5
6	2022	6

```
In [20]: data = data.merge(idx, on='year')
```

## Models - Predict 2020

### Data Sets

```
In [21]: data_train = data[data['year'] < 2020]
data_test = data[data['year'].isin([2020])]
```

```
In [22]: data_train.sort_values(['country', 'year'])[1:5]
```

```
Out[22]:
```

	year	iso	country	region	EPI	AIR	H2O	BDH	WRS	idx
615	2012	ALB	Albania	Eastern Europe	65.85	100.00	72.28	61.72	23.19	1
492	2014	ALB	Albania	Eastern Europe	54.73	68.24	55.91	63.19	3.36	2
369	2016	ALB	Albania	Eastern Europe	29.61	37.41	48.46	26.07	2.70	3
246	2018	ALB	Albania	Eastern Europe	38.80	37.10	51.50	57.10	2.70	4

```
In [23]: data_test[1:5]
```

```
Out[23]:
```

	year	iso	country	region	EPI	AIR	H2O	BDH	WRS	idx
124	2020	DZA	Algeria	Greater Middle East	44.8	45.3	53.2	39.0	33.1	5
125	2020	AGO	Angola	Sub-Saharan Africa	29.7	26.8	12.8	39.3	0.0	5
126	2020	ARG	Argentina	Latin America & Caribbean	52.2	56.9	64.7	49.1	5.9	5
127	2020	ARM	Armenia	Former Soviet States	52.3	36.3	57.2	79.2	8.8	5

## Simple Linear Regression

```
In [24]: # Train-predict and plot for all countries
x_train = data_train[['idx']]
y_train = data_train['EPI']

x_test = data_test[['idx']]
y_test = data_test['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + p.const
data_plot_train['tag'] = 'train'
data_plot_test = data_test
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + p.const
data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=True)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

#### OLS Regression Results

```
=====
Dep. Variable:          EPI      R-squared:                0.099
Model:                  OLS      Adj. R-squared:           0.098
Method:                 Least Squares      F-statistic:         67.46
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):    1.28e-15
Time:                  03:18:21      Log-Likelihood:       -2506.3
No. Observations:      615      AIC:                  5017.
Df Residuals:          613      BIC:                  5025.
Df Model:              1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	59.0999	0.996	59.313	0.000	57.143	61.057
idx	-3.3410	0.407	-8.213	0.000	-4.140	-2.542

```
=====
Omnibus:                23.974      Durbin-Watson:           1.812
Prob(Omnibus):          0.000      Jarque-Bera (JB):        13.442
Skew:                   0.189      Prob(JB):                0.00121
Kurtosis:               2.382      Cond. No.                 4.74
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykerne
l_launcher.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stabl
e/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykerne
l_launcher.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stabl
e/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykerne
l_launcher.py:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stabl
e/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykerne
l_launcher.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stabl
e/user_guide/indexing.html#returning-a-view-versus-a-copy
Out[24]: <AxesSubplot:xlabel='idx', ylabel='fitted'>

```



```

In [25]: x_test = sm.add_constant(x_test, has_constant='add')
         predictions = model.predict(x_test)

```

```

In [26]: mt.mean_squared_error(y_test, predictions)

```

```

Out[26]: 320.56405438271065

```

Sample Country

```
In [27]: # Train-predict and plot for sample country Norway
country = 'Norway'
x_train = data_train[data_train['country'] == country][['idx']]
y_train = data_train[data_train['country'] == country]['EPI']

x_test = data_test[data_test['country'] == country][['idx']]
y_test = data_test[data_test['country'] == country]['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train[data_train['country'] == country]
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + p.const
data_plot_train['tag'] = 'train'
data_plot_test = data_test[data_test['country'] == country]
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + p.const
data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=True)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\stats\stattools.py:75: ValueWarning: omni\_normtest is not valid with less than 8 observations; 5 samples were given.

"samples were given." % int(n), ValueWarning)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:20: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:22: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:23: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

## OLS Regression Results

Dep. Variable:	EPI	R-squared:	0.429
Model:	OLS	Adj. R-squared:	0.239
Method:	Least Squares	F-statistic:	2.256
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	0.230
Time:	03:18:42	Log-Likelihood:	-13.755
No. Observations:	5	AIC:	31.51
Df Residuals:	3	BIC:	30.73
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	78.2340	3.789	20.650	0.000	66.177	90.291
idx	-2.3230	1.547	-1.502	0.230	-7.245	2.599


  

Omnibus:	nan	Durbin-Watson:	3.484
Prob(Omnibus):	nan	Jarque-Bera (JB):	0.484
Skew:	-0.430	Prob(JB):	0.785
Kurtosis:	1.741	Cond. No.	4.74

### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Out[27]: <AxesSubplot:xlabel='idx', ylabel='fitted'>



In [28]: `x_test = sm.add_constant(x_test, has_constant='add')`  
`predictions = model.predict(x_test)`

In [29]: `mt.mean_squared_error(y_test, predictions)`


Out[29]: 122.78856100000007

## Multiple Linear Regression

In [30]: `sns.regplot(x='AIR', y='EPI', data=data)`



Out[30]: <AxesSubplot:xlabel='AIR', ylabel='EPI'>



```
In [31]: x_train = data_train[['idx', 'AIR', 'H2O', 'BDH', 'WRS']]
y_train = data_train['EPI']

x_test = data_test[['idx', 'AIR', 'H2O', 'BDH', 'WRS']]
y_test = data_test['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + \
                             data_plot_train['AIR'] * p.AIR + \
                             data_plot_train['H2O'] * p.H2O + \
                             data_plot_train['BDH'] * p.BDH + \
                             data_plot_train['WRS'] * p.WRS + p.const

data_plot_train['tag'] = 'train'
data_plot_test = data_test
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + \
                             data_plot_test['AIR'] * p.AIR + \
                             data_plot_test['H2O'] * p.H2O + \
                             data_plot_test['BDH'] * p.BDH + \
                             data_plot_test['WRS'] * p.WRS + p.const

data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=False)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

# OLS Regression Results

```

=====
Dep. Variable:          EPI      R-squared:                0.846
Model:                  OLS      Adj. R-squared:           0.845
Method:                  Least Squares      F-statistic:           668.4
Date:                   Sun, 11 Dec 2022      Prob (F-statistic):       1.62e-244
Time:                   03:19:03      Log-Likelihood:          -1963.4
No. Observations:       615      AIC:                    3939.
Df Residuals:           609      BIC:                    3965.
Df Model:               5
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	17.7819	0.956	18.610	0.000	15.905	19.658
idx	0.1390	0.189	0.734	0.463	-0.233	0.511
AIR	0.1306	0.012	10.521	0.000	0.106	0.155
H2O	0.1854	0.012	15.749	0.000	0.162	0.209
BDH	0.1582	0.010	15.343	0.000	0.138	0.178
WRS	0.1622	0.010	16.698	0.000	0.143	0.181

```

=====
Omnibus:                6.288      Durbin-Watson:           1.926
Prob(Omnibus):          0.043      Jarque-Bera (JB):        6.337
Skew:                   -0.198      Prob(JB):                0.0421
Kurtosis:               3.300      Cond. No.                 475.
=====

```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:21: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:22: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:28: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:29: SettingWithCopyWarning:


A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

<AxesSubplot:xlabel='idx', ylabel='fitted'>

Out[31]:



```
In [32]: fig = sm.graphics.plot_partregress_grid(model)
fig.tight_layout(pad=1.0)
```

```
eval_env: 1
eval_env: 1
eval_env: 1
eval_env: 1
eval_env: 1
eval_env: 1
```



```
In [33]: x_test = sm.add_constant(x_test, has_constant='add')
predictions = model.predict(x_test)
```

```
In [34]: mt.mean_squared_error(y_test, predictions)
```

```
Out[34]: 20.794376334029756
```

## Sample Country

```
In [35]: # Train-predict and plot for sample country Norway
x_train = data_train[data_train['country'] == country][['idx', 'AIR', 'H2O', 'BDH']
```

```

y_train = data_train[data_train['country'] == country]['EPI']

x_test = data_test[data_test['country'] == country][['idx', 'AIR', 'H2O', 'BDH', 'WRS']]
y_test = data_test[data_test['country'] == country]['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train[data_train['country'] == country]
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + \
    data_plot_train['AIR'] * p.AIR + \
    data_plot_train['H2O'] * p.H2O + \
    data_plot_train['BDH'] * p.BDH + \
    data_plot_train['WRS'] * p.WRS + p.const

data_plot_train['tag'] = 'train'
data_plot_test = data_test[data_test['country'] == country]
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + \
    data_plot_test['AIR'] * p.AIR + \
    data_plot_test['H2O'] * p.H2O + \
    data_plot_test['BDH'] * p.BDH + \
    data_plot_test['WRS'] * p.WRS + p.const

data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=False)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)

```

```
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\stats\stattools.py:75: ValueWarning: omni_normtest is not valid with less than 8 observations; 5 samples were given.
"samples were given." % int(n), ValueWarning)
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear_model.py:1765: RuntimeWarning: divide by zero encountered in true_divide
    return 1 - (np.divide(self.nobs - self.k_constant, self.df_resid)
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear_model.py:1766: RuntimeWarning: invalid value encountered in double_scalars
    * (1 - self.rsquared))
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear_model.py:1687: RuntimeWarning: divide by zero encountered in double_scalars
    return np.dot(wresid, wresid) / self.df_resid
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:23: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:29: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:30: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

# OLS Regression Results

```

=====
Dep. Variable:          EPI    R-squared:                1.000
Model:                  OLS    Adj. R-squared:             nan
Method:                 Least Squares    F-statistic:             nan
Date:                  Sun, 11 Dec 2022    Prob (F-statistic):         nan
Time:                  03:19:23    Log-Likelihood:            144.69
No. Observations:      5    AIC:                      -279.4
Df Residuals:          0    BIC:                      -281.3
Df Model:              4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0034	inf	0	nan	nan	nan
idx	-2.0882	inf	-0	nan	nan	nan
AIR	-0.3160	inf	-0	nan	nan	nan
H2O	0.6633	inf	0	nan	nan	nan
BDH	0.3145	inf	0	nan	nan	nan
WRS	0.3167	inf	0	nan	nan	nan

```

=====
Omnibus:                nan    Durbin-Watson:            0.660
Prob(Omnibus):           nan    Jarque-Bera (JB):         0.359
Skew:                   -0.607    Prob(JB):                 0.836
Kurtosis:                2.500    Cond. No.                 748.
=====


```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The input rank is higher than the number of observations.

Out[35]: <AxesSubplot:xlabel='idx', ylabel='fitted'>



```

In [36]: x_test = sm.add_constant(x_test, has_constant='add')
         predictions = model.predict(x_test)

```

```

In [37]: mt.mean_squared_error(y_test, predictions)

```

Out[37]: 97.68311940770529

# Models - Predict 2022

## Data Sets

```
In [38]: data_train = data[(data['year'] < 2022)]
data_test = data[data['year'].isin([2022])]
```

## Simple Linear Regression

```
In [39]: x_train = data_train[['idx']]
y_train = data_train['EPI']

x_test = data_test[['idx']]
y_test = data_test['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + p.const
data_plot_train['tag'] = 'train'
data_plot_test = data_test
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + p.const
data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=True)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

# OLS Regression Results

```

=====
Dep. Variable:          EPI      R-squared:          0.057
Model:                  OLS      Adj. R-squared:       0.056
Method:                 Least Squares      F-statistic:       44.67
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):  4.61e-11
Time:                  03:19:40      Log-Likelihood:     -3030.8
No. Observations:      738      AIC:                6066.
Df Residuals:          736      BIC:                6075.
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	57.4727	0.961	59.827	0.000	55.587	59.359
idx	-2.1206	0.317	-6.684	0.000	-2.744	-1.498

```

=====
Omnibus:              30.015      Durbin-Watson:       1.770
Prob(Omnibus):        0.000      Jarque-Bera (JB):     16.693
Skew:                 0.197      Prob(JB):             0.000237
Kurtosis:             2.377      Cond. No.             5.78
=====

```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
<AxesSubplot:xlabel='idx', ylabel='fitted'>
```

Out[39]:





```
In [40]: x_test = sm.add_constant(x_test, has_constant='add')
         predictions = model.predict(x_test)
```

```
In [41]: mt.mean_squared_error(y_test, predictions)
```

```
Out[41]: 183.67397634532637
```

## Sample Country

```
In [42]: # Train-predict and plot for sample country Norway
         x_train = data_train[data_train['country'] == country][['idx']]
         y_train = data_train[data_train['country'] == country]['EPI']

         x_test = data_test[data_test['country'] == country][['idx']]
         y_test = data_test[data_test['country'] == country]['EPI']

         # with statsmodels
         x_train = sm.add_constant(x_train) # adding a constant
         model = sm.OLS(y_train, x_train).fit()
         predictions_train = model.predict(x_train)

         print_model = model.summary()
         print(print_model)

         p = model.params
         data_plot_train = data_train[data_train['country'] == country]
         data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + p.const
         data_plot_train['tag'] = 'train'
         data_plot_test = data_test[data_test['country'] == country]
         data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + p.const
         data_plot_test['tag'] = 'test'
         data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

         sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=True)
         sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\stats\stattools.py:75: ValueWarning: omni_normtest is not valid with less than 8 observations; 6 samples were given.
"samples were given." % int(n), ValueWarning)
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:19: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

#### OLS Regression Results

```


=====
Dep. Variable:          EPI      R-squared:                0.069
Model:                  OLS      Adj. R-squared:           -0.164
Method:                 Least Squares      F-statistic:          0.2943
Date:                   Sun, 11 Dec 2022    Prob (F-statistic):      0.616
Time:                   03:20:47           Log-Likelihood:        -17.746
No. Observations:       6              AIC:                   39.49
Df Residuals:           4              BIC:                   39.08
Df Model:               1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          76.1233      4.130      18.433      0.000      64.657      87.589
idx           -0.7400      1.364      -0.543      0.616     -4.527      3.047
=====
Omnibus:            nan      Durbin-Watson:           2.573
Prob(Omnibus):      nan      Jarque-Bera (JB):         0.866
Skew:              -0.028     Prob(JB):              0.648
Kurtosis:           1.139     Cond. No.               5.78
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Out[42]: <AxesSubplot:xlabel='idx', ylabel='fitted'>



```
In [43]: x_test = sm.add_constant(x_test, has_constant='add')
         predictions = model.predict(x_test)
```

```
In [44]: mt.mean_squared_error(y_test, predictions)
```

```
Out[44]: 153.3469444444439
```

## Multiple Linear Regression

```
In [45]: x_train = data_train[['idx', 'AIR', 'H2O', 'BDH', 'WRS']]
         y_train = data_train['EPI']

         x_test = data_test[['idx', 'AIR', 'H2O', 'BDH', 'WRS']]
         y_test = data_test['EPI']

         # with statsmodels
         x_train = sm.add_constant(x_train) # adding a constant
         model = sm.OLS(y_train, x_train).fit()
         predictions_train = model.predict(x_train)

         print_model = model.summary()
         print(print_model)

         p = model.params
         data_plot_train = data_train
         data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + \
                                     data_plot_train['AIR'] * p.AIR + \
                                     data_plot_train['H2O'] * p.H2O + \
                                     data_plot_train['BDH'] * p.BDH + \
                                     data_plot_train['WRS'] * p.WRS + p.const

         data_plot_train['tag'] = 'train'
         data_plot_test = data_test
         data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + \
                                    data_plot_test['AIR'] * p.AIR + \
                                    data_plot_test['H2O'] * p.H2O + \
                                    data_plot_test['BDH'] * p.BDH + \
                                    data_plot_test['WRS'] * p.WRS + p.const

         data_plot_test['tag'] = 'test'
```

```
data_plot = pd.concat([data_plot_train,data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=False)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)
```

#### OLS Regression Results

```
=====
Dep. Variable:          EPI      R-squared:                0.859
Model:                  OLS      Adj. R-squared:            0.859
Method:                 Least Squares      F-statistic:        895.4
Date:                  Sun, 11 Dec 2022      Prob (F-statistic):    5.09e-309
Time:                  03:20:58      Log-Likelihood:       -2328.5
No. Observations:      738      AIC:                  4669.
Df Residuals:          732      BIC:                  4697.
Df Model:               5
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	16.8606	0.824	20.470	0.000	15.244	18.478
idx	0.2858	0.137	2.082	0.038	0.016	0.555
AIR	0.1320	0.011	11.569	0.000	0.110	0.154
H2O	0.1968	0.011	18.136	0.000	0.176	0.218
BDH	0.1626	0.009	17.305	0.000	0.144	0.181
WRS	0.1545	0.009	18.049	0.000	0.138	0.171

```
=====
Omnibus:                9.169      Durbin-Watson:          1.941
Prob(Omnibus):          0.010      Jarque-Bera (JB):       10.622
Skew:                   -0.179      Prob(JB):               0.00494
Kurtosis:               3.466      Cond. No.                459.
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:21: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:22: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:28: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:29: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Out[45]:



In [46]:

In [47]:

Out[47]:

## Sample Country

In [48]:

```
# Train-predict and plot for sample country Norway
x_train = data_train[data_train['country'] == country][['idx', 'AIR', 'H2O', 'BDH']]
y_train = data_train[data_train['country'] == country]['EPI']

x_test = data_test[data_test['country'] == country][['idx', 'AIR', 'H2O', 'BDH']]
y_test = data_test[data_test['country'] == country]['EPI']

# with statsmodels
x_train = sm.add_constant(x_train) # adding a constant
model = sm.OLS(y_train, x_train).fit()
predictions_train = model.predict(x_train)

print_model = model.summary()
print(print_model)

p = model.params
data_plot_train = data_train[data_train['country'] == country]
data_plot_train['fitted'] = data_plot_train['idx'] * p.idx + \
    data_plot_train['AIR'] * p.AIR + \
    data_plot_train['H2O'] * p.H2O + \
    data_plot_train['BDH'] * p.BDH + \
    data_plot_train['WRS'] * p.WRS + p.const
data_plot_train['tag'] = 'train'
data_plot_test = data_test[data_test['country'] == country]
data_plot_test['fitted'] = data_plot_test['idx'] * p.idx + \
    data_plot_test['AIR'] * p.AIR + \
    data_plot_test['H2O'] * p.H2O + \
    data_plot_test['BDH'] * p.BDH + \
    data_plot_test['WRS'] * p.WRS + p.const
data_plot_test['tag'] = 'test'
```

```

data_plot_test['WRS'] * p.WRS + p.const
data_plot_test['tag'] = 'test'
data_plot = pd.concat([data_plot_train, data_plot_test], axis=0)

sns.lmplot(x='idx', y='EPI', data=data_plot, hue='tag', fit_reg=False)
sns.regplot(x='idx', y='fitted', data=data_plot, color='red', scatter=False)

```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\stats\stattools.py:75: ValueWarning: omni\_normtest is not valid with less than 8 observations; 6 samples were given.

"samples were given." % int(n), ValueWarning)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1765: RuntimeWarning: divide by zero encountered in true\_divide

return 1 - (np.divide(self.nobs - self.k\_constant, self.df\_resid)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1766: RuntimeWarning: invalid value encountered in double\_scalars

\* (1 - self.rsquared))

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1687: RuntimeWarning: divide by zero encountered in double\_scalars

return np.dot(wresid, wresid) / self.df\_resid

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:22: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:23: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:29: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:30: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

# OLS Regression Results

```

=====
Dep. Variable:          EPI    R-squared:                1.000
Model:                OLS    Adj. R-squared:            nan
Method:              Least Squares    F-statistic:            nan
Date:                Sun, 11 Dec 2022    Prob (F-statistic):      nan
Time:                03:21:30    Log-Likelihood:         146.03
No. Observations:      6    AIC:                    -280.1
Df Residuals:          0    BIC:                    -281.3
Df Model:              5
Covariance Type:      nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1390.5510	inf	-0	nan	nan	nan
idx	0.5853	inf	0	nan	nan	nan
AIR	-1.3900	inf	-0	nan	nan	nan
H2O	15.8391	inf	0	nan	nan	nan
BDH	0.0554	inf	0	nan	nan	nan
WRS	0.2105	inf	0	nan	nan	nan

```

=====
Omnibus:                nan    Durbin-Watson:            0.002
Prob(Omnibus):          nan    Jarque-Bera (JB):         0.205
Skew:                   0.446    Prob(JB):                 0.903
Kurtosis:               2.838    Cond. No.                 2.91e+05
=====


```

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.91e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Out[48]: <AxesSubplot:xlabel='idx', ylabel='fitted'>



```

In [49]: x_test = sm.add_constant(x_test, has_constant='add')
         predictions = model.predict(x_test)

```

```

In [50]: mt.mean_squared_error(y_test, predictions)

```

Out[50]: 708.281056503761

## ARIMA

### Data prep

```
In [51]: d = {'year': [2010, 2012, 2014, 2016, 2018, 2020, 2022],  
            'date': ['2010-01-01', '2012-01-01', '2014-01-01', '2016-01-01', '2018-01-01', '2020-01-01', '2022-01-01'],  
            'data_lookup': [60.251138, 71.842927, 80.631463, 57.126748, 61.361057, 32.226748, 2.0]}  
data_lookup = pd.DataFrame(data=d)
```

```
In [52]: data = data.merge(data_lookup, on='year')
```

### Predict 2020

### All Countries Avg

```
In [53]: data_arima_all = data.groupby(['date', 'year'], as_index=False).mean()
```

```
In [54]: data_arima_all
```

```
Out[54]:
```

	date	year	EPI	AIR	H2O	BDH	WRS	idx
0	2010-01-01	2010	60.251138	67.837967	77.512683	58.312114	68.656992	0.0
1	2012-01-01	2012	53.542439	71.842927	57.831789	60.059106	31.310000	1.0
2	2014-01-01	2014	55.251382	80.631463	57.126748	61.361057	32.226748	2.0
3	2016-01-01	2016	45.454146	43.607154	46.896911	53.589024	30.661789	3.0
4	2018-01-01	2018	47.590244	44.652033	51.738211	58.256098	30.661789	4.0
5	2020-01-01	2020	50.937398	49.315447	55.434146	62.035772	30.661789	5.0
6	2022-01-01	2022	44.358537	44.569106	55.507317	58.218699	32.001626	6.0

```
In [55]: data_train_arima = data_arima_all[data_arima_all['year'] < 2020].sort_values('year')  
data_test_arima = data_arima_all[data_arima_all['year'].isin([2020])].sort_values('year')
```

```
In [56]: data_train_arima
```

```
Out[56]:
```

	date	year	EPI	AIR	H2O	BDH	WRS	idx
0	2010-01-01	2010	60.251138	67.837967	77.512683	58.312114	68.656992	0.0
1	2012-01-01	2012	53.542439	71.842927	57.831789	60.059106	31.310000	1.0
2	2014-01-01	2014	55.251382	80.631463	57.126748	61.361057	32.226748	2.0
3	2016-01-01	2016	45.454146	43.607154	46.896911	53.589024	30.661789	3.0
4	2018-01-01	2018	47.590244	44.652033	51.738211	58.256098	30.661789	4.0

```
In [57]: data_train_arima = data_train_arima[['date', 'EPI']]  
data_test_arima = data_test_arima[['date', 'EPI']]  
  
data_train_arima.set_index('date', inplace=True)  
data_test_arima.set_index('date', inplace=True)
```



```

data_train_arima.index = pd.to_datetime(data_train_arima.index)
data_test_arima.index = pd.to_datetime(data_test_arima.index)

stepwise_model = auto_arima(data_train_arima,
                             start_p=1, start_q=1,
                             max_p=5, max_q=5,
                             m=1,
                             seasonal=False,
                             d=1, trace=True,
                             error_action='ignore',
                             suppress_warnings=True,
                             stepwise=False)

print(stepwise_model.aic())

stepwise_model.fit(data_train_arima)

future_forecast = stepwise_model.predict(n_periods=1)
data_test_arima['Prediction'] = future_forecast
data_arima = pd.concat([data_train_arima, data_test_arima])
data_arima['idx'] = data_arima.index #.astype(str)

```

```

ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=28.550, Time=0.03 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=inf, Time=0.08 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=22.223, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=inf, Time=0.12 sec

```

```

Best model: ARIMA(1,1,0)(0,0,0)[0] intercept
Total fit time: 0.265 seconds
22.22342320426996

```

In [58]: future\_forecast

Out[58]: 2020-01-01 38.791233  
Freq: 2AS-JAN, dtype: float64

In [59]: data\_arima[['EPI', 'Prediction']]

Out[59]:


	EPI	Prediction
2010-01-01	60.251138	NaN
2012-01-01	53.542439	NaN
2014-01-01	55.251382	NaN
2016-01-01	45.454146	NaN
2018-01-01	47.590244	NaN
2020-01-01	50.937398	38.791233

In [60]:

```

plt.plot(data_arima.idx.dt.year, data_arima.EPI, 'o-', label = "Actual")
plt.scatter(data_arima.idx.dt.year, data_arima.Prediction, label = "Prediction", c=
plt.legend()
plt.ylim([25, 70])
plt.show()

```



```
In [61]: mt.mean_squared_error(data_test_arima.EPI, data_test_arima.Prediction)
```

```
Out[61]: 147.52932490439358
```

## Sample Country

```
In [62]: data_train_arima = data[data['year'] < 2020][data['country'] == country].sort_values('year')
data_test_arima = data[data['year'].isin([2020])[data['country'] == country].sort_values('year')]
```

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

"""Entry point for launching an IPython kernel.

C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel\_launcher.py:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
In [63]: data_train_arima = data_train_arima[['date', 'EPI']]
data_test_arima = data_test_arima[['date', 'EPI']]

data_train_arima.set_index('date', inplace=True)
data_test_arima.set_index('date', inplace=True)

data_train_arima.index = pd.to_datetime(data_train_arima.index)
data_test_arima.index = pd.to_datetime(data_test_arima.index)
```

```
In [64]: data_train_arima
```

```
Out[64]:
```

	EPI
date	
2010-01-01	81.13
2012-01-01	69.92
2014-01-01	78.04
2016-01-01	68.75
2018-01-01	70.10

```
In [65]: data_test_arima
```

Out[65]: **EPI**

	date
2020-01-01	77.7

```
In [66]: stepwise_model = auto_arma(data_train_arma,
                                   start_p=1, start_q=1,
                                   max_p=5, max_q=5,
                                   m=1,
                                   seasonal=False,
                                   d=1, trace=True,
                                   error_action='ignore',
                                   suppress_warnings=True,
                                   stepwise=False)

print(stepwise_model.aic())

ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=31.881, Time=0.01 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=inf, Time=0.07 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=27.418, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=inf, Time=0.09 sec

Best model: ARIMA(1,1,0)(0,0,0)[0] intercept
Total fit time: 0.197 seconds
27.417919922228585
```

```
In [67]: stepwise_model.fit(data_train_arma)
```

```
Out[67]: ARIMA(maxiter=50, method='lbfgs', order=(1, 1, 0), out_of_sample_size=0,
               scoring='mse', scoring_args={}, seasonal_order=(0, 0, 0, 0),
               start_params=None, suppress_warnings=True, trend=None,
               with_intercept=True)
```


```
In [68]: future_forecast = stepwise_model.predict(n_periods=1)
data_test_arma['Prediction'] = future_forecast
data_arma = pd.concat([data_train_arma, data_test_arma])
data_arma['idx'] = data_arma.index#.astype(str)
```

```
In [69]: data_arma[['EPI', 'Prediction']]
```

Out[69]: **EPI Prediction**

	date	
2010-01-01	81.13	NaN
2012-01-01	69.92	NaN
2014-01-01	78.04	NaN
2016-01-01	68.75	NaN
2018-01-01	70.10	NaN
2020-01-01	77.70	64.872379

```
In [70]: plt.plot(data_arma.idx.dt.year, data_arma.EPI, 'o-', label = "Actual")
plt.scatter(data_arma.idx.dt.year, data_arma.Prediction, label = "Prediction", c=
plt.legend()
plt.show()
```



## Predict 2022

## All Countries

```
In [71]: data_train_arma = data_arma_all[data_arma_all['year'] < 2022].sort_values('year')
data_test_arma = data_arma_all[data_arma_all['year'].isin([2022])].sort_values('year')
```

```
In [72]: data_train_arma = data_train_arma[['date', 'EPI']]
data_test_arma = data_test_arma[['date', 'EPI']]

data_train_arma.set_index('date', inplace=True)
data_test_arma.set_index('date', inplace=True)

data_train_arma.index = pd.to_datetime(data_train_arma.index)
data_test_arma.index = pd.to_datetime(data_test_arma.index)

stepwise_model = auto_arma(data_train_arma,
                           start_p=1, start_q=1,
                           max_p=5, max_q=5,
                           m=1,
                           seasonal=False,
                           d=1, trace=True,
                           error_action='ignore',
                           suppress_warnings=True,
                           stepwise=False)

print(stepwise_model.aic())

stepwise_model.fit(data_train_arma)

future_forecast = stepwise_model.predict(n_periods=1)
data_test_arma['Prediction'] = future_forecast
data_arma = pd.concat([data_train_arma, data_test_arma])
data_arma['idx'] = data_arma.index #.astype(str)
```

```

ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=34.933, Time=0.01 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=34.852, Time=0.10 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.10 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=35.868, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=2438603.608, Time=0.05 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.11 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=39.553, Time=0.14 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=38.514, Time=0.15 sec
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=inf, Time=0.23 sec

```

Best model: ARIMA(0,1,1)(0,0,0)[0] intercept  
Total fit time: 0.921 seconds  
34.8522281264441


In [73]: `data_arma[['EPI', 'Prediction']]`

Out[73]:

	EPI	Prediction
date		

date		
2010-01-01	60.251138	NaN
2012-01-01	53.542439	NaN
2014-01-01	55.251382	NaN
2016-01-01	45.454146	NaN
2018-01-01	47.590244	NaN
2020-01-01	50.937398	NaN
2022-01-01	44.358537	44.756014

In [74]: `plt.plot(data_arma.idx, data_arma.EPI, 'o-', label = "Actual")`  
`plt.scatter(data_arma.idx, data_arma.Prediction, label = "Prediction", color='red')`  
`plt.legend()`  
`plt.ylim([25, 70])`  
`plt.show()`



In [75]: `mt.mean_squared_error(data_test_arma.EPI, data_test_arma.Prediction)`

Out[75]: 0.15798856585889806

Sample Country

```
In [76]: data_train_arima = data[data['year'] < 2022][data['country'] == country].sort_values()
data_test_arima = data[data['year'].isin([2022])[data['country'] == country].sort_values()
```

```
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
"""Entry point for launching an IPython kernel.
C:\Users\ozlem.gunes\AppData\Local\Continuum\miniconda3\lib\site-packages\ipykernel_launcher.py:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.
```

```
In [77]: data_train_arima = data_train_arima[['date', 'EPI']]
data_test_arima = data_test_arima[['date', 'EPI']]

data_train_arima.set_index('date', inplace=True)
data_test_arima.set_index('date', inplace=True)

data_train_arima.index = pd.to_datetime(data_train_arima.index)
data_test_arima.index = pd.to_datetime(data_test_arima.index)
```

```
In [78]: data_train_arima
```

```
Out[78]:
```

	EPI
date	
2010-01-01	81.13
2012-01-01	69.92
2014-01-01	78.04
2016-01-01	68.75
2018-01-01	70.10
2020-01-01	77.70

```
In [79]: data_test_arima
```

```
Out[79]:
```

	EPI
date	
2022-01-01	59.3

```
In [80]: stepwise_model = auto_arima(data_train_arima,
                                     start_p=1, start_q=1,
                                     max_p=5, max_q=5,
                                     m=1,
                                     seasonal=False,
                                     d=1, trace=True,
                                     error_action='ignore',
                                     suppress_warnings=True,
                                     stepwise=False)

print(stepwise_model.aic())
```

```

ARIMA(0,1,0)(0,0,0)[0] intercept      : AIC=39.215, Time=0.01 sec
ARIMA(0,1,1)(0,0,0)[0] intercept      : AIC=inf, Time=0.13 sec
ARIMA(0,1,2)(0,0,0)[0] intercept      : AIC=inf, Time=0.10 sec
ARIMA(1,1,0)(0,0,0)[0] intercept      : AIC=38.598, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept      : AIC=inf, Time=0.12 sec
ARIMA(1,1,2)(0,0,0)[0] intercept      : AIC=inf, Time=0.13 sec
ARIMA(2,1,0)(0,0,0)[0] intercept      : AIC=37.529, Time=0.12 sec
ARIMA(2,1,1)(0,0,0)[0] intercept      : AIC=1397421.593, Time=0.05 sec
ARIMA(2,1,2)(0,0,0)[0] intercept      : AIC=inf, Time=0.19 sec

```

```

Best model:  ARIMA(2,1,0)(0,0,0)[0] intercept
Total fit time: 0.877 seconds
37.528581428326504

```

```
In [81]: stepwise_model.fit(data_train_arma)
```

```
Out[81]: ARIMA(maxiter=50, method='lbfgs', order=(2, 1, 0), out_of_sample_size=0,
            scoring='mse', scoring_args={}, seasonal_order=(0, 0, 0, 0),
            start_params=None, suppress_warnings=True, trend=None,
            with_intercept=True)
```


```
In [82]: future_forecast = stepwise_model.predict(n_periods=1)
data_test_arma['Prediction'] = future_forecast
data_arma = pd.concat([data_train_arma, data_test_arma])
data_arma['idx'] = data_arma.index.astype(str)
```

```
In [83]: data_arma[['EPI', 'Prediction']]
```

```
Out[83]:
```

	EPI Prediction	
date		
2010-01-01	81.13	NaN
2012-01-01	69.92	NaN
2014-01-01	78.04	NaN
2016-01-01	68.75	NaN
2018-01-01	70.10	NaN
2020-01-01	77.70	NaN
2022-01-01	59.30	61.724263

```
In [84]: plt.plot(data_arma.idx, data_arma.EPI, 'o-', label = "Actual")
plt.scatter(data_arma.idx, data_arma.Prediction, label = "Prediction", color='red')
plt.legend()
plt.show()
```



In [ ]: